

Operating Systems

Project 2: Paging Simulation

Osniel Quintana
Felipe Gutierrez

INTRODUCTION:

Page replacement is an important component of an operating system. When a page containing a desired piece of information or instruction is searched in translation buffers or page tables and found missing from main memory, a page fault is said to occur. A page fault is a type of interrupt, raised by the hardware when a running program accesses a memory page that is mapped into the virtual address space, but not loaded in physical memory. As the size of main memory is limited and is much smaller than the size of physical memory, the role of page replacement is to identify the best page to remove from main memory as a result of a page fault and replace it by the new page from disk that contains the requested value or instruction. The page replacement algorithm used in the project are First-In-First-Out, Least Recently Used , and the VMS Second Chance Algorithm.

This project uses a simulator to conduct FIFO, LRU or VMS page replacement policies when accessing memory specified by the traces given for this project. The simulator maintains a Page Table, as well as it's insertions and removals. When the Page Table is full, and a new address needs to be loaded, one of the replacements policies are used to decide which page to remove. When pages are loaded into memory (into the Page Table), the queue pertaining to FIFO, LRU or VMS is updated accordingly. Similarly, when a page hit occurs (an address is already in memory), the page is referenced for LRU or VMS.

The user has control over the number of frames used in the page table and the trace file to be used for the simulation. These parameters can be passed as arguments on terminal when the simulator is to be executed.

The purpose of this project is to determine the overall efficiency of these three different replacement policies when applied to large trace files (containing 1 million entries). Multiple

frames sizes will also be used in the calculation of this efficiency. The measures used to calculate the overall performance of these policies will be the total number of reads from disk and writes to disk and the hit rate, given some frames that the user inputs.

Time Spent: 48 hours

METHODS:

The algorithms we used in this project as the replacement policies are FIFO, LRU and VMS Second Chance Algorithm. FIFO keeps a list of all pages currently in memory, the page at the beginning of the list being the oldest one and the page at the end being the most recent one. When the Page Table is full, and some space needs to be made for another page, FIFO replacement policy evicts the oldest page in memory (first one to arrive). Least Recently Used policy replaces the least-frequently used page when an eviction must take place. Pages are added to the queue as they are loaded into memory, but when a page is referenced and it is in memory, the page is brought to the front of the queue. LRU evicts pages the same way FIFO does but maintains its queue differently, as previously shown. The last method that we used is the VMS second chance lists where pages are placed in two different process queues when loaded into memory and instead of being evicted from memory they are loaded into a global clean-page free list and dirty-page list, from which they might be evicted later on. When a process queue is out of space, it removes one of its pages in FIFO order and adds it to the global clean or dirty queue depending on whether the page was dirty or not (if it was ever written to). If another process Q needs a free page, it takes the first free page off of the global clean list. However, if a page is referenced from memory, and it is present in the global queues, the appropriate process reclaims it (enqueues) in its queue. The bigger these global second-chance lists are, the closer the VMS algorithm performs to LRU.

For the experiments we ran all three replacement policies in all four traces using frames sizes from 4 to 4096 (or 512 for bzip.trace) which corresponds to cache size of 16 KB to 16384 KB (16 MB). We increment the frame size in powers of 2 for two reasons: to make it easier to calculate the cache size needed and because increments of 1 or some other random number seemed inefficient. As the results will show, choosing these increments clearly shows when the maximum efficiency is achieved. To calculate the cache size, we multiply the number of frames by 4 since each frame (page) is 4 KB in size per the assignment instructions. The results were recorded and the annexed graphs show the hit rate vs cache size and the amount of reads and writes of each replacement policy per trace file. The hit rate is calculated by the formula:

$$hit\ rate = \frac{1000000 - (reads + writes)}{1000000} * 100$$

Where *reads* is the number of reads from disk and *writes* is the number of writes to disk. The hit rate is displayed as a percentage, and the size of each trace file was shown to be of 1 million entries. This hit rate is hence the number of page hits, a page hit is anything that is not a read or a write to disk. Every time a page is loaded into memory the *reads* increase, and when a page is evicted from memory and it was ever written the *writes* to disk increase.

To improve upon the efficiency of our LRU replacement algorithm, we implemented a HashMap that maintains a mapping from page number of the table to a pointer of the place in the LRU queue where this page is located. The LRU queue is implemented on a double linked list, and this in combination with the hash map made it trivial to move a node on the queue to the front while updating the rest of the pointers. The other two replacement policies are implemented in a standard way also using the double linked list based queue.

RESULTS:

After testing the three algorithms of choice we notice that as the frame size increases the disk access (reads or writes) decreases drastically. This is because having fewer frames will increase the number of page faults because of the lower freedom in replacement choice, in other words, the algorithms don't have enough pages to choose from. On the other hand, having too many pages becomes expensive in terms of memory.

Analyzing the Hit rate vs Cache Size, we can see that except for the bzip.trace, all other traces need at least 2 MB of memory to maximize the hit rate. Most contemporary memories are in the size of 8 GB (8192 MB) but taking into consideration how many other processes use memory and the fact that a single process that handles this trace needs 2 MB of memory, it becomes evident that this size of memory is a little excessive. All algorithms in all traces perform poorly with low cache size (< 64 KB), but as the cache size increases the algorithms need to replace less pages, hence more hits are achieved. The bzip.trace appeared to be an exception. This trace achieves its maximum hit rate at 512 KB under all replacement policies. We believe this is partly due to the nature of the traces it contains: probably the virtual page number of those addresses are actually a limited number, meaning that although there is one million traces, most of the traces reference the same page multiple times for a read or write. Due to this nature, the maximum hit rate is easier to achieve as the replacement policies won't be called as often to evict a page.

Overall the LRU algorithm seemed to perform the best. It achieves a higher Hit Rate than all other policies with lower cache size. However, in some traces as in bzip.trace and gcc.trace, VMS was slightly better than LRU. VMS is supposed to come really close to the performance of LRU as the cache size goes up, but the fact that it outperforms LRU in some traces indicates that

these traces contains addresses that are referenced really often. Since VMS maintains two additional global queues (clean and dirty) to keep pages that should be evicted, it reduces the number of evictions as long as these pages are referenced multiple times. LRU will evict a page before VMS even if the page is referenced too often, as LRU does not give a “second chance” to the page and waits for it to be referenced again. For these reasons, LRU performs really well on swim.trace and sixpack.trace. Our intuition is that these files have a larger number of unique VPN, which decreases the chances of a hit and increases the number of evictions for a new page.

FIFO has the overall poorest performance (as expected) since it evicts the first page to arrive when a replacement is needed but does not account for a page that was recently referenced. With FIFO, pages in the middle of the queue might be referenced but will be treated as “old” eventually and be evicted, when in reality the chances of the page being referenced again increased.


CONCLUSION:



Throughout this project we have used various page replacement algorithms like FIFO, LRU, and VMS Second Chance then simulated and compared them to evaluate their efficiency. Upon analyzing the results, it becomes evident that LRU performs best in most situations. The situations that favor LRU are those where not a lot of page hits will happen even if the memory was big enough to fit the memory space. VMS favors from the opposite, having traces that reference the same page number multiple times. FIFO will achieve a performance similar to VMS and LRU if the pages being referenced maintain the same FIFO order as in the FIFO queue, in other words, if pages that are referenced after being loaded are currently at the end of the queue (the newest in the queue).



Based on the number of reads and writes for each replacement policy, the bzip.trace seems to have the higher chances of referencing the same page, whereas the swim.trace and sixpack.trace contain a lot of unique virtual page numbers that are not referenced as often. Therefore, LRU performed really well on swim.trace and sixpack.trace, but VMS outperformed LRU on bzip.trace. The higher the memory space (cache size), the higher the hit rate all algorithms achieved. For all traces, 2 MB was enough to achieve near 100% hit rate, and 512 KB was enough to achieve 100% in the case of bzip.trace.

We have seen that out of all three replacement policies, FIFO performs the worst. VMS, depending on the situation, can outperform or come close to LRU, as indicated in the hit rate vs cache size for all traces. Since we can't know the nature of the traces before execution, we concluded that the safest bet is to use LRU as it achieved a high hit rate on average.

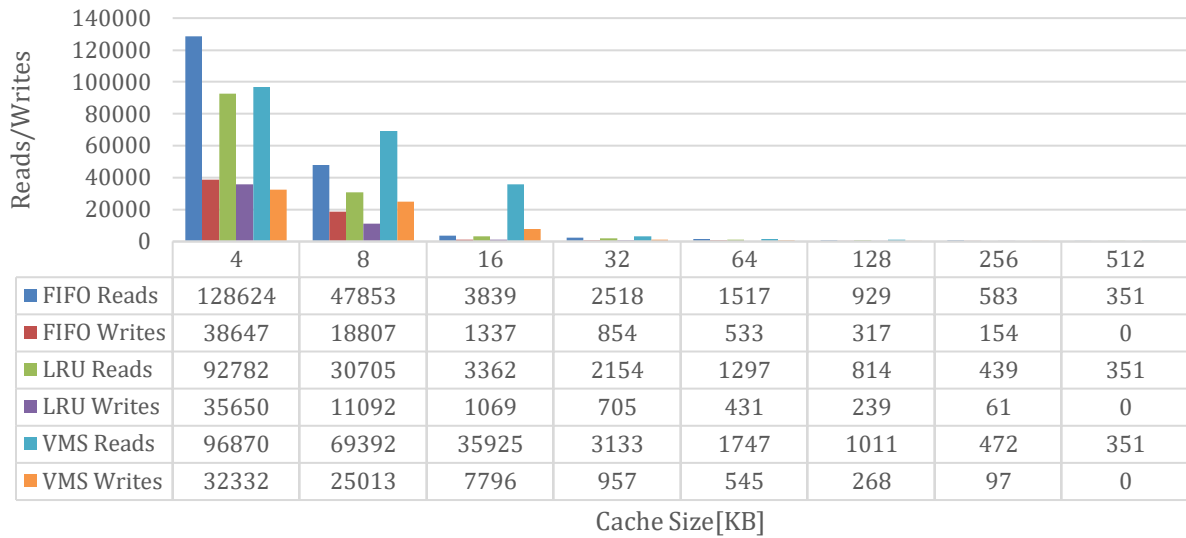
ANNEX:







bzip.trace Reads/Writes vs Frames



sixpack.trace Reads/Writes vs Frames

